

# object detection

March 9, 2022

## 1 Image Classification and Object Localization

```
[ ]: # In this lab, you'll build a CNN from scratch to:  
  
# classify the main subject in an image  
# localize it by drawing bounding boxes around it.
```

```
[1]: ##Import  
import os, re, time, json  
import PIL.Image, PIL.ImageFont, PIL.ImageDraw  
import numpy as np  
try:  
    # %tensorflow_version only exists in Colab.  
    %tensorflow_version 2.x  
except Exception:  
    pass  
import tensorflow as tf  
from matplotlib import pyplot as plt  
import tensorflow_datasets as tfds  
  
print("Tensorflow version " + tf.__version__)
```

Using TensorFlow backend.

```
-----  
ModuleNotFoundError                                Traceback (most recent call last)  
~\AppData\Local\Temp\ipykernel_22080\1621094636.py in <module>  
    10 import tensorflow as tf  
    11 from matplotlib import pyplot as plt  
----> 12 import tensorflow_datasets as tfds  
    13  
    14 print("Tensorflow version " + tf.__version__)  
  
ModuleNotFoundError: No module named 'tensorflow_datasets'
```

```
[4]: #Data visualization
```

```
#@title Plot Utilities for Bounding Boxes [RUN ME]
```

```
im_width = 75  
im_height = 75
```

```
use_normalized_coordinates = True
```

```
def draw_bounding_boxes_on_image_array(image,  
                                       boxes,  
                                       color=[],  
                                       thickness=1,  
                                       display_str_list=()):  
  
    """Draws bounding boxes on image (numpy array).  
    Args:  
        image: a numpy array object.  
        boxes: a 2 dimensional numpy array of [N, 4]: (ymin, xmin, ymax, xmax).  
                The coordinates are in normalized format between [0, 1].  
        color: color to draw bounding box. Default is red.  
        thickness: line thickness. Default value is 4.  
        display_str_list_list: a list of strings for each bounding box.  
    Raises:  
        ValueError: if boxes is not a [N, 4] array  
    """  
  
    image_pil = PIL.Image.fromarray(image)  
    rgbimg = PIL.Image.new("RGBA", image_pil.size)  
    rgbimg.paste(image_pil)  
    draw_bounding_boxes_on_image(rgbimg, boxes, color, thickness,  
                                display_str_list)  
  
    return np.array(rgbimg)
```

```
def draw_bounding_boxes_on_image(image,  
                                 boxes,  
                                 color=[],  
                                 thickness=1,  
                                 display_str_list=()):  
  
    """Draws bounding boxes on image.  
    Args:  
        image: a PIL.Image object.  
        boxes: a 2 dimensional numpy array of [N, 4]: (ymin, xmin, ymax, xmax).  
                The coordinates are in normalized format between [0, 1].  
        color: color to draw bounding box. Default is red.  
        thickness: line thickness. Default value is 4.  
        display_str_list: a list of strings for each bounding box.
```

```

Raises:
    ValueError: if boxes is not a [N, 4] array
    """
boxes_shape = boxes.shape
if not boxes_shape:
    return
if len(boxes_shape) != 2 or boxes_shape[1] != 4:
    raise ValueError('Input must be of size [N, 4]')
for i in range(boxes_shape[0]):
    draw_bounding_box_on_image(image, boxes[i, 1], boxes[i, 0], boxes[i, 3],
                              boxes[i, 2], color[i], thickness,
→display_str_list[i])

def draw_bounding_box_on_image(image,
                              ymin,
                              xmin,
                              ymax,
                              xmax,
                              color='red',
                              thickness=1,
                              display_str=None,
                              use_normalized_coordinates=True):
    """Adds a bounding box to an image.
    Bounding box coordinates can be specified in either absolute (pixel) or
    normalized coordinates by setting the use_normalized_coordinates argument.
    Args:
        image: a PIL.Image object.
        ymin: ymin of bounding box.
        xmin: xmin of bounding box.
        ymax: ymax of bounding box.
        xmax: xmax of bounding box.
        color: color to draw bounding box. Default is red.
        thickness: line thickness. Default value is 4.
        display_str_list: string to display in box
        use_normalized_coordinates: If True (default), treat coordinates
            ymin, xmin, ymax, xmax as relative to the image. Otherwise treat
            coordinates as absolute.
    """
    draw = PIL.ImageDraw.Draw(image)
    im_width, im_height = image.size
    if use_normalized_coordinates:
        (left, right, top, bottom) = (xmin * im_width, xmax * im_width,
                                       ymin * im_height, ymax * im_height)
    else:
        (left, right, top, bottom) = (xmin, xmax, ymin, ymax)
    draw.line([(left, top), (left, bottom), (right, bottom),
              (right, top), (left, top)], width=thickness, fill=color)

```

```
[5]: """
This cell contains helper functions used for visualization
and downloads only.

You can skip reading it, as there is very
little Keras or Tensorflow related code here.
"""

# Matplotlib config

plt.rc('image', cmap='gray')
plt.rc('grid', linewidth=0)
plt.rc('xtick', top=False, bottom=False, labelsizе='large')
plt.rc('ytick', left=False, right=False, labelsizе='large')
plt.rc('axes', facecolor='F8F8F8', titlesize="large", edgecolor='white')
plt.rc('text', color='a8151a')
plt.rc('figure', facecolor='FOFOFO') # Matplotlib fonts
MATPLOTLIB_FONT_DIR = os.path.join(os.path.dirname(plt.__file__), "mpl-data/
→fonts/ttf")

# pull a batch from the datasets. This code is not very nice, it gets much
→better in eager mode (TODO)
def dataset_to_numpy_util(training_dataset, validation_dataset, N):

    # get one batch from each: 10000 validation digits, N training digits
    batch_train_ds = training_dataset.unbatch().batch(N)

    # eager execution: loop through datasets normally
    if tf.executing_eagerly():
        for validation_digits, (validation_labels, validation_bboxes) in
→validation_dataset:
            validation_digits = validation_digits.numpy()
            validation_labels = validation_labels.numpy()
            validation_bboxes = validation_bboxes.numpy()
            break
        for training_digits, (training_labels, training_bboxes) in batch_train_ds:
            training_digits = training_digits.numpy()
            training_labels = training_labels.numpy()
            training_bboxes = training_bboxes.numpy()
            break

    # these were one-hot encoded in the dataset
    validation_labels = np.argmax(validation_labels, axis=1)
    training_labels = np.argmax(training_labels, axis=1)
```

```

    return (training_digits, training_labels, training_bboxes,
            validation_digits, validation_labels, validation_bboxes)

# create digits from local fonts for testing
def create_digits_from_local_fonts(n):
    font_labels = []
    img = PIL.Image.new('LA', (75*n, 75), color = (0,255)) # format 'LA': black
    ↪ in channel 0, alpha in channel 1
    font1 = PIL.ImageFont.truetype(os.path.join(MATPLOTLIB_FONT_DIR,
    ↪ 'DejaVuSansMono-Oblique.ttf'), 25)
    font2 = PIL.ImageFont.truetype(os.path.join(MATPLOTLIB_FONT_DIR, 'STIXGeneral.
    ↪ ttf'), 25)
    d = PIL.ImageDraw.Draw(img)
    for i in range(n):
        font_labels.append(i%10)
        d.text((7+i*75, 0 if i<10 else -4), str(i%10), fill=(255,255), font=font1 if
    ↪ i<10 else font2)
    font_digits = np.array(img.getdata(), np.float32)[: ,0] / 255.0 # black in
    ↪ channel 0, alpha in channel 1 (discarded)
    font_digits = np.reshape(np.stack(np.split(np.reshape(font_digits, [75,
    ↪ 75*n]), n, axis=1), axis=0), [n, 75*75])
    return font_digits, font_labels

# utility to display a row of digits with their predictions
def display_digits_with_boxes(digits, predictions, labels, pred_bboxes, bboxes,
    ↪ iou, title):

    n = 10

    indexes = np.random.choice(len(predictions), size=n)
    n_digits = digits[indexes]
    n_predictions = predictions[indexes]
    n_labels = labels[indexes]

    n_iou = []
    if len(iou) > 0:
        n_iou = iou[indexes]

    if (len(pred_bboxes) > 0):
        n_pred_bboxes = pred_bboxes[indexes,:]

    if (len(bboxes) > 0):
        n_bboxes = bboxes[indexes,:]

```

```

n_digits = n_digits * 255.0
n_digits = n_digits.reshape(n, 75, 75)
fig = plt.figure(figsize=(20, 4))
plt.title(title)
plt.yticks([])
plt.xticks([])

for i in range(10):
    ax = fig.add_subplot(1, 10, i+1)
    bboxes_to_plot = []
    if (len(pred_bboxes) > i):
        bboxes_to_plot.append(n_pred_bboxes[i])

    if (len(bboxes) > i):
        bboxes_to_plot.append(n_bboxes[i])

    img_to_draw = draw_bounding_boxes_on_image_array(image=n_digits[i],
→boxes=np.asarray(bboxes_to_plot), color=['red', 'green'],
→display_str_list=["true", "pred"])
    plt.xlabel(n_predictions[i])
    plt.xticks([])
    plt.yticks([])

    if n_predictions[i] != n_labels[i]:
        ax.xaxis.label.set_color('red')

plt.imshow(img_to_draw)

if len(iou) > i :
    color = "black"
    if (n_iou[i][0] < iou_threshold):
        color = "red"
    ax.text(0.2, -0.3, "iou: %s" %(n_iou[i][0]), color=color, transform=ax.
→transAxes)

# utility to display training and validation curves
def plot_metrics(metric_name, title, ylim=5):
    plt.title(title)
    plt.ylim(0,ylim)
    plt.plot(history.history[metric_name],color='blue',label=metric_name)
    plt.plot(history.history['val_' + metric_name],color='green',label='val_' +
→metric_name)

```

```
[ ]: # There will be a training folder and a testing folder.
# Each of these will have a subfolder for cats and another subfolder for dogs.
```

```
[6]: try:
    os.mkdir('/tmp/cats-v-dogs')
    os.mkdir('/tmp/cats-v-dogs/training')
    os.mkdir('/tmp/cats-v-dogs/testing')
    os.mkdir('/tmp/cats-v-dogs/training/cats')
    os.mkdir('/tmp/cats-v-dogs/training/dogs')
    os.mkdir('/tmp/cats-v-dogs/testing/cats')
    os.mkdir('/tmp/cats-v-dogs/testing/dogs')
except OSError:
    pass
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>  
170500096/170498071 [=====] - 55s 0us/step  
170508288/170498071 [=====] - 55s 0us/step

```
[1]: # Detect hardware
try:
    tpu = tf.distribute.cluster_resolver.TPUClusterResolver() # TPU detection
except ValueError:
    tpu = None
    gpus = tf.config.experimental.list_logical_devices("GPU")

# Select appropriate distribution strategy
if tpu:
    tf.config.experimental_connect_to_cluster(tpu)
    tf.tpu.experimental.initialize_tpu_system(tpu)
    strategy = tf.distribute.experimental.TPUStrategy(tpu) # Going back and forth
    ↪ between TPU and host is expensive. Better to run 128 batches on the TPU
    ↪ before reporting back.
    print('Running on TPU ', tpu.cluster_spec().as_dict()['worker'])
elif len(gpus) > 1:
    strategy = tf.distribute.MirroredStrategy([gpu.name for gpu in gpus])
    print('Running on multiple GPUs ', [gpu.name for gpu in gpus])
elif len(gpus) == 1:
    strategy = tf.distribute.get_strategy() # default strategy that works on CPU
    ↪ and single GPU
    print('Running on single GPU ', gpus[0].name)
else:
    strategy = tf.distribute.get_strategy() # default strategy that works on CPU
    ↪ and single GPU
    print('Running on CPU')
print("Number of accelerators: ", strategy.num_replicas_in_sync)
```

```
[2]: BATCH_SIZE = 64 * strategy.num_replicas_in_sync # Goba batch size.
# The global batch size will be automatically sharded across all
# replicas by the tf.data.Dataset API. A single TPU has 8 cores.
# The best practice is to scale the batch size by the numb
```

```
-----
NameError                                Traceback (most recent call last)
~\AppData\Local\Temp\ipykernel_37072\3703507671.py in <module>
    35
    36 split_size = .9
----> 37 split_data(CAT_SOURCE_DIR, TRAINING_CATS_DIR, TESTING_CATS_DIR,
    ↪split_size)
    38 split_data(DOG_SOURCE_DIR, TRAINING_DOGS_DIR, TESTING_DOGS_DIR,
    ↪split_size)
    39

~\AppData\Local\Temp\ipykernel_37072\3703507671.py in split_data(SOURCE,
    ↪TRAINING, TESTING, SPLIT_SIZE)
      3 def split_data(SOURCE, TRAINING, TESTING, SPLIT_SIZE):
      4     files = []
----> 5     for filename in os.listdir(SOURCE):
      6         file = SOURCE + filename
      7         if os.path.getsize(file) > 0:

NameError: name 'os' is not defined
```

```
[9]: '''
Transforms each image in dataset by pasting it on a 75x75 canvas at random
    ↪locations.
'''
def read_image_tfds(image, label):
    xmin = tf.random.uniform((), 0 , 48, dtype=tf.int32)
    ymin = tf.random.uniform((), 0 , 48, dtype=tf.int32)
    image = tf.reshape(image, (28,28,1))
    image = tf.image.pad_to_bounding_box(image, ymin, xmin, 75, 75)
    image = tf.cast(image, tf.float32)/255.0
    xmin = tf.cast(xmin, tf.float32)
    ymin = tf.cast(ymin, tf.float32)

    xmax = (xmin + 28) / 75
    ymax = (ymin + 28) / 75
    xmin = xmin / 75
    ymin = ymin / 75
    return image, (tf.one_hot(label, 10), [xmin, ymin, xmax, ymax])

'''
```



*Loads and maps the training split of the dataset using the map function. Note*  
*→ that we try to load the gcs version since TPU can only work with datasets on*  
*→ Google Cloud Storage.*

*'''*

```
def get_training_dataset():
```

```
    with strategy.scope():
```

```
        dataset = tfds.load("mnist", split="train", as_supervised=True,
```

```
        → try_gcs=True)
```

```
        dataset = dataset.map(read_image_tfds, num_parallel_calls=16)
```

```
        dataset = dataset.shuffle(5000, reshuffle_each_iteration=True)
```

```
        dataset = dataset.repeat() # Mandatory for Keras for now
```

```
        dataset = dataset.batch(BATCH_SIZE, drop_remainder=True) #
```

```
        → drop_remainder is important on TPU, batch size must be fixed
```

```
        dataset = dataset.prefetch(-1) # fetch next batches while training on
```

```
        → the current one (-1: autotune prefetch buffer size)
```

```
        return dataset
```

*'''*

*Loads and maps the validation split of the dataset using the map function. Note*  
*→ that we try to load the gcs version since TPU can only work with datasets on*  
*→ Google Cloud Storage.*

*'''*

```
def get_validation_dataset():
```

```
    dataset = tfds.load("mnist", split="test", as_supervised=True, try_gcs=True)
```

```
    dataset = dataset.map(read_image_tfds, num_parallel_calls=16)
```

```
    #dataset = dataset.cache() # this small dataset can be entirely cached in
```

```
    → RAM
```

```
    dataset = dataset.batch(10000, drop_remainder=True) # 10000 items in eval
```

```
    → dataset, all in one batch
```

```
    dataset = dataset.repeat() # Mandatory for Keras for now
```

```
    return dataset
```

```
# instantiate the datasets
```

```
with strategy.scope():
```

```
    training_dataset = get_training_dataset()
```

```
    validation_dataset = get_validation_dataset()
```

## 2 Visualize Data

```
[10]: (training_digits, training_labels, training_bboxes,  
       validation_digits, validation_labels, validation_bboxes) =  
       → dataset_to_numpy_util(training_dataset, validation_dataset, 10)
```

```
display_digits_with_boxes(training_digits, training_labels, training_labels, np.
    ↳array([]), training_bboxes, np.array([]), "training digits and their labels")
display_digits_with_boxes(validation_digits, validation_labels,
    ↳validation_labels, np.array([]), validation_bboxes, np.array([]),
    ↳"validation digits and their labels")
```

[11]: *# Define the Network*

```
# Here, you'll define your custom CNN.

# feature_extractor: these convolutional layers extract the features of the
    ↳image.
# classifier: This define the output layer that predicts among 10 categories
    ↳(digits 0 through 9)
# bounding_box_regression: This defines the output layer that predicts 4
    ↳numeric values, which define the coordinates of the bounding box (xmin,
    ↳ymin, xmax, ymax)
# final_model: This combines the layers for feature extraction, classification
    ↳and bounding box prediction.
# Notice that this is another example of a branching model, because the model
    ↳splits to produce two kinds of output (a category and set of numbers).
# Since you've learned to use the Functional API earlier in the specialization
    ↳(course 1), you have the flexibility to define this kind of branching model!
# define_and_compile_model: choose the optimizer and metrics, then compile the
    ↳model.
```

[ ]: *# Get and prepare the model*  
*# You'll be using the InceptionV3 model.*

```
# Since you're making use of transfer learning, you'll load the pre-trained
    ↳weights of the model.
# You'll also freeze the existing layers so that they aren't trained on your
    ↳downstream task with the cats and dogs data.
# You'll also get a reference to the last layer, 'mixed7' because you'll add
    ↳some layers after this last layer.
```

[12]: *'''*  
*Feature extractor is the CNN that is made up of convolution and pooling layers.*  
*'''*

```
def feature_extractor(inputs):
    x = tf.keras.layers.Conv2D(16, activation='relu', kernel_size=3,
    ↳input_shape=(75, 75, 1))(inputs)
    x = tf.keras.layers.AveragePooling2D((2, 2))(x)

    x = tf.keras.layers.Conv2D(32, kernel_size=3, activation='relu')(x)
```

```

x = tf.keras.layers.AveragePooling2D((2, 2))(x)

x = tf.keras.layers.Conv2D(64, kernel_size=3, activation='relu')(x)
x = tf.keras.layers.AveragePooling2D((2, 2))(x)

    return x

'''
dense_layers adds a flatten and dense layer.
This will follow the feature extraction layers
'''
def dense_layers(inputs):
    x = tf.keras.layers.Flatten()(inputs)
    x = tf.keras.layers.Dense(128, activation='relu')(x)
    return x

'''
Classifier defines the classification output.
This has a set of fully connected layers and a softmax layer.
'''
def classifier(inputs):

    classification_output = tf.keras.layers.Dense(10, activation='softmax', name='classification')(inputs)
    return classification_output

'''
This function defines the regression output for bounding box prediction.
Note that we have four outputs corresponding to (xmin, ymin, xmax, ymax)
'''
def bounding_box_regression(inputs):
    bounding_box_regression_output = tf.keras.layers.Dense(units = '4', name = 'bounding_box')(inputs)
    return bounding_box_regression_output

def final_model(inputs):
    feature_cnn = feature_extractor(inputs)
    dense_output = dense_layers(feature_cnn)

    '''
    The model branches here.
    The dense layer's output gets fed into two branches:
    classification_output and bounding_box_output
    '''

```

```

        classification_output = classifier(dense_output)
        bounding_box_output = bounding_box_regression(dense_output)

        model = tf.keras.Model(inputs = inputs, outputs = [classification_output,
↪ bounding_box_output])

        return model

def define_and_compile_model(inputs):
    model = final_model(inputs)

    model.compile(optimizer='adam',
                  loss = {'classification' : 'categorical_crossentropy',
                          'bounding_box' : 'mse'
                          },
                  metrics = {'classification' : 'accuracy',
                             'bounding_box' : 'mse'
                             })

    return model

with strategy.scope():
    inputs = tf.keras.layers.Input(shape=(75, 75, 1,))
    model = define_and_compile_model(inputs)

# print model layers
model.summary()

```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 32, 32, 3)]	0
up_sampling2d (UpSampling2D)	(None, 224, 224, 3)	0
resnet50 (Functional)	(None, 7, 7, 2048)	23587712
global_average_pooling2d (G1	(None, 2048)	0
flatten (Flatten)	(None, 2048)	0
dense (Dense)	(None, 1024)	2098176
dense_1 (Dense)	(None, 512)	524800
classification (Dense)	(None, 10)	5130

```

=====
Total params: 26,215,818
Trainable params: 26,162,698
Non-trainable params: 53,120
-----

```

```

[ ]: # Train and validate the model
EPOCHS = 10 # 45
steps_per_epoch = 60000//BATCH_SIZE # 60,000 items in this dataset
validation_steps = 1

history = model.fit(training_dataset,
                    steps_per_epoch=steps_per_epoch,
                    ↪validation_data=validation_dataset, validation_steps=validation_steps,
                    ↪epochs=EPOCHS)

loss, classification_loss, bounding_box_loss, classification_accuracy,
↪bounding_box_mse = model.evaluate(validation_dataset, steps=1)
print("Validation accuracy: ", classification_accuracy)

```

### 3 Intersection over union

```

[ ]: def intersection_over_union(pred_box, true_box):
    xmin_pred, ymin_pred, xmax_pred, ymax_pred = np.split(pred_box, 4, axis =
    ↪1)
    xmin_true, ymin_true, xmax_true, ymax_true = np.split(true_box, 4, axis = 1)

    smoothing_factor = 1e-10

    xmin_overlap = np.maximum(xmin_pred, xmin_true)
    xmax_overlap = np.minimum(xmax_pred, xmax_true)
    ymin_overlap = np.maximum(ymin_pred, ymin_true)
    ymax_overlap = np.minimum(ymax_pred, ymax_true)

    pred_box_area = (xmax_pred - xmin_pred) * (ymax_pred - ymin_pred)
    true_box_area = (xmax_true - xmin_true) * (ymax_true - ymin_true)

    overlap_area = np.maximum((xmax_overlap - xmin_overlap), 0) * np.
    ↪maximum((ymax_overlap - ymin_overlap), 0)
    union_area = (pred_box_area + true_box_area) - overlap_area

    iou = (overlap_area + smoothing_factor) / (union_area + smoothing_factor)

    return iou

```

## 4 Visualize predictions

The following code will make predictions and visualize both the classification and the predicted bounding boxes.

The true bounding box labels will be in green, and the model's predicted bounding boxes are in red. The predicted number is shown below the image.

```
[ ]: # recognize validation digits
predictions = model.predict(validation_digits, batch_size=64)
predicted_labels = np.argmax(predictions[0], axis=1)

predicted_bboxes = predictions[1]

iou = intersection_over_union(predicted_bboxes, validation_bboxes)

iou_threshold = 0.6

print("Number of predictions where iou > threshold(%s): %s" % (iou_threshold,
    ↳(iou >= iou_threshold).sum()))
print("Number of predictions where iou < threshold(%s): %s" % (iou_threshold,
    ↳(iou < iou_threshold).sum()))

display_digits_with_boxes(validation_digits, predicted_labels,
    ↳validation_labels, predicted_bboxes, validation_bboxes, iou, "True and
    ↳Predicted values")
```